

Problem Statement and Goals

Software Engineering

Team 6, Blockchainerz
Robert Lam
Nathan Hum
Aidan Lao
Keira Laskoski

Table 1: Revision History

Date	Developer(s)	Change
September 24	Robert	Write problem statement and inputs/outputs
Date2	Name(s)	Description of changes
...	...	...

1 Problem Statement

[You should check your problem statement with the [problem statement checklist](#) —SS].

[You can change the section headings, as long as you include the required information. —SS]

1.1 Problem

Current AI inference markets are increasingly concentrated around a small number of large organizations [?]. As the artificial intelligence (AI) industry continues to grow, we become increasingly reliant on centralized providers for AI models, data, hardware, and compute. This centralization increases the risk of biased outputs because of the lack of diversity in training algorithms and data used by these companies [?]. Furthermore, the oligopolized nature of the AI market creates a risk of censorship and manipulation of AI outputs. The lack of transparency from these companies and the black-box nature of many AI models makes it difficult for consumers to verify the quality and reliability of their inferences. Lastly, the centralized nature of the compute infrastructure

used by these companies creates potential single points of failure. According to the 2026 Stanford AI Index, NVIDIA accounts for over 60% of total global AI compute capacity [?]. This concentration poses a significant risk of disruption to AI markets in the event of a single failure.

[What problem are you trying to solve? Why is it important? Do not talk about how you will solve the problem, only what the problem is. AI is rarely part of the problem. The problem is “what” you want to do, not “how” you want to do it. The problem statement should be written in a way that is understandable to a non-technical audience. —SS]

1.2 Inputs and Outputs

The inputs to the system consists of task definitions, inference requests, and network participation. Task definitions are specification of a problem to be solved via machine learning, including the expected input and output types. Inference requests ask the network to perform inference for a given task and include input data or some specification of where to find the input data. Network participation includes machine learning predictions for individual inference requests that are supplied to the network by an individual.

The system outputs aggregated inference results and maintains information about the network. The aggregated inference results include the final prediction for a given inference request, as well as information about the quality of the prediction.

[Characterize the problem in terms of “high level” inputs and outputs. Use abstraction so that you can avoid details. —SS]

1.3 Stakeholders

- **Task requesters**

Task requesters are people who are requesting inferences to be made for a given task. These people would value the decentralized nature of our platform. For example, healthcare researchers would make requests to predict various regression inferences (such as biomarker levels) while still complying with privacy constraints. Similarly, decentralized finance platforms could make requests to predict various regression inferences (such as stock prices) while still complying with privacy constraints.

- **Inference providers**

Inference providers are people with computing resources and models, and want to monetize their regression models. These people usually have specialized or powerful models that can be used to make predictions for a given task.

- **Data Providers**

Data providers are people with data that can be used to train models for

a given task. They would be interested in helping the inference providers improve their models by providing them with data to train on.

1.4 Environment

This software is compatible with any computer running a modern browser that is connected to wifi. There is no hardware component.

2 Goals

This project has a total of five primary goals.

2.1 Automated and Decentralized Marketplace Transactions

Enable users and model providers to autonomously establish and fulfill marketplace agreements without requiring centralized coordination. This goal will be measured by whether the system can successfully complete an agreement from the initial request through payment and settlement without manual intervention.

2.2 Accessible AI Model Marketplace

Provide users with a marketplace where AI models can be discovered, evaluated, and selected based on factors such as model reputation and availability. This goal will be measured by whether users can discover and select available models through the marketplace without needing to interact directly with the underlying decentralized infrastructure.

2.3 Reliable AI Inference Through the Marketplace

Enable users to submit AI inference requests through the marketplace and receive results from participating model providers. This goal will be measured by whether a submitted inference request can successfully travel from the consumer to a model provider and return a corresponding inference result.

2.4 Accessible Consumer Experience

Provide a consumer-facing experience that allows users to interact with the marketplace without requiring knowledge of decentralized systems or their underlying infrastructure. This goal will be measured by whether a user can discover a model, submit an inference request, and receive its result entirely through the consumer interface.

2.5 Model Provider Participation

Enable independent model providers to participate in the marketplace by making their models available to consumers and responding to inference requests. This goal will be measured by whether a model provider can register a model, receive an inference request, execute the inference, and return the result through the system.

3 Stretch Goals

This project has a total of two stretch goals, which outline additional features that could be implemented after the primary goals have been achieved.

3.1 Expanded AI Model and Inference Support

Extend the marketplace beyond linear regression to support additional AI models and more complex inference tasks. This is considered a stretch goal because supporting multiple model architectures and more complex inference workflows would require additional development beyond the core marketplace functionality established by the primary goals. The achievement of this goal would be measured by whether the system can successfully register, discover, and execute inference requests for multiple AI model types beyond linear regression.

3.2 Private AI Inference

Enable inference on private or sensitive data while protecting that data from unauthorized access during execution through the use of Trusted Execution Environments (TEEs). This is considered a stretch goal because incorporating TEEs would introduce additional security, infrastructure, and integration requirements beyond those necessary to support the core marketplace and inference functionality. The achievement of this goal would be measured by whether the system can successfully execute an inference request on private data within a TEE while preventing unauthorized access to the data during execution.

[Goals you hope to get to, but not necessary. They are also helpful if the scope of the original project needs to be expanded. —SS]

4 Custom Documents (CDs)

[For CAS 741: State whether the project is a research project. This designation, with the approval (or request) of the instructor, can be modified over the course of the term. —SS]

[For SE Capstone: List your custom documents (CDs). Potential CDs include usability testing, code walkthroughs, user documentation, formal proof, GenderMag personas, Design Thinking, etc. (The full list is on the course outline and in Lecture 02.) Normally the number of CDs will be two (one per term).

Approval of the CDs will be part of the discussion with the instructor for approving the project. The CDs, with the approval (or request) of the instructor, can be modified over the course of the term. —SS]

Appendix — Reflection

[Not required for CAS 741 —SS]

The purpose of reflection questions is to give you a chance to assess your own learning and that of your group as a whole, and to find ways to improve in the future. Reflection is an important part of the learning process. Reflection is also an essential component of a successful software development process.

Reflections are most interesting and useful when they're honest, even if the stories they tell are imperfect. You will be marked based on your depth of thought and analysis, and not based on the content of the reflections themselves. Thus, for full marks we encourage you to answer openly and honestly and to avoid simply writing "what you think the evaluator wants to hear."

Please answer the following questions. Some questions can be answered on the team level, but where appropriate, each team member should write their own response:

1. What went well while writing this deliverable?
2. What pain points did you experience during this deliverable, and how did you resolve them?
3. How did you and your team adjust the scope of your goals to ensure they are suitable for a Capstone project (not overly ambitious but also of appropriate complexity for a senior design project)?